

# Healthcare Analytics Project Implementation

April 11, 2023

```
[1]: # Niranjan J 19MIA1003  
     # Roshan Srinivaas S 19MIA1001
```

```
[61]: import numpy as np  
import pandas as pd  
import os  
import cv2 as cv  
import matplotlib.pyplot as plt  
import seaborn as sns
```

# Load and Check Data

```
[62]: all_0 = "../input/leukemia-classification/C-NMC_Leukemia/training_data/fold_0/  
        ↪all"  
all_1 = "../input/leukemia-classification/C-NMC_Leukemia/training_data/fold_1/  
        ↪all"  
all_2 = "../input/leukemia-classification/C-NMC_Leukemia/training_data/fold_2/  
        ↪all"  
  
hem_0 = "../input/leukemia-classification/C-NMC_Leukemia/training_data/fold_0/  
        ↪hem"  
hem_1 = "../input/leukemia-classification/C-NMC_Leukemia/training_data/fold_1/  
        ↪hem"  
hem_2 = "../input/leukemia-classification/C-NMC_Leukemia/training_data/fold_2/  
        ↪hem"
```

```
[63]: def get_path_image(folder):  
    image_paths = []  
    image_fnames = os.listdir(folder)  
    for img_id in range(len(image_fnames)):  
        img = os.path.join(folder, image_fnames[img_id])  
        image_paths.append(img)  
  
    return image_paths
```

```
[64]: img_data = []
```

```

for i in [all_0,all_1,all_2,hem_0,hem_1,hem_2]:
    paths = get_path_image(i)
    img_data.extend(paths)
print(len(img_data))

```

10661

```

[65]: data = {"img_data":img_data,
            "labels":[np.nan for x in range(len(img_data))]}

data = pd.DataFrame(data)

```

```

[66]: data["labels"][0:7272] = 1 # ALL
data["labels"][7272:10661] = 0 # HEM

```

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:1:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

"""Entry point for launching an IPython kernel.

/opt/conda/lib/python3.7/site-packages/ipykernel\_launcher.py:2:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy)

```

[67]: data["labels"] = data["labels"].astype("int64")


```

# Crop Black Edges In Image

```

[68]: image = cv.imread(data["img_data"][1000])
plt.imshow(image)
plt.title("Sample image before cropping")
plt.show()


```



```
[69]: img_list = []
      for i in range(len(img_data)):
          image = cv.imread(data["img_data"][i])
          gray = cv.cvtColor(image, cv.COLOR_BGR2GRAY)
          thresh = cv.threshold(gray, 0, 255, cv.THRESH_BINARY_INV + cv.
          ↪THRESH_OTSU)[1]

          result = cv.bitwise_and(image, image, mask=thresh)
          result[result==0] = [255,255,255]
          (x, y, z_) = np.where(result > 0)
          mnx = (np.min(x))
          maxx = (np.max(x))
          mny = (np.min(y))
          mxy = (np.max(y))
          crop_img = image[mnx:maxx,mny:mxy,:]
          crop_img_r = cv.resize(crop_img, (224,224))
          img_list.append(crop_img_r)
```

```
[70]: plt.imshow(img_list[1000])
      plt.title("Sample image after cropping")
      plt.show()
```



# Feature Extraction with ResNet50

```
[71]: from tensorflow.keras.applications import ResNet50, ResNet101
      from keras.applications.vgg19 import VGG19
      from tensorflow.keras.preprocessing import image
      from tensorflow.keras.models import Model
      from tensorflow.keras.applications.resnet50 import preprocess_input
```

```
[72]: def feature_extract(model):
      if model == "VGG19": model = VGG19(weights='imagenet', include_top=False,
      ↪pooling="avg")
      elif model == "ResNet50": model =
      ↪ResNet50(weights='imagenet', include_top=False, pooling="avg")
      elif model == "ResNet101": model =
      ↪ResNet101(weights='imagenet', include_top=False, pooling="avg")
      return model
```

```
[73]: model = feature_extract("ResNet50")
```

```
[74]: features_list = []
      for i in range(len(img_list)):
```

```

image = img_list[i].reshape(-1, 224, 224, 3)
image = preprocess_input(image)
features = model.predict(image).reshape(2048, )
features_list.append(features)

```

```

1/1 [=====] - 1s 1s/step
1/1 [=====] - 0s 134ms/step
1/1 [=====] - 0s 134ms/step
1/1 [=====] - 0s 143ms/step
1/1 [=====] - 0s 148ms/step
1/1 [=====] - 0s 126ms/step
1/1 [=====] - 0s 127ms/step
1/1 [=====] - 0s 126ms/step
1/1 [=====] - 0s 153ms/step
1/1 [=====] - 0s 135ms/step
1/1 [=====] - 0s 143ms/step
1/1 [=====] - 0s 134ms/step
1/1 [=====] - 0s 132ms/step
1/1 [=====] - 0s 141ms/step
1/1 [=====] - 0s 126ms/step
1/1 [=====] - 0s 137ms/step
1/1 [=====] - 0s 134ms/step
1/1 [=====] - 0s 127ms/step
1/1 [=====] - 0s 139ms/step
1/1 [=====] - 0s 128ms/step
1/1 [=====] - 0s 127ms/step
1/1 [=====] - 0s 141ms/step
1/1 [=====] - 0s 144ms/step
1/1 [=====] - 0s 137ms/step
1/1 [=====] - 0s 135ms/step
1/1 [=====] - 0s 147ms/step
1/1 [=====] - 0s 129ms/step
1/1 [=====] - 0s 138ms/step
1/1 [=====] - 0s 137ms/step
1/1 [=====] - 0s 135ms/step
1/1 [=====] - 0s 146ms/step
1/1 [=====] - 0s 144ms/step
1/1 [=====] - 0s 133ms/step
1/1 [=====] - 0s 133ms/step
1/1 [=====] - 0s 130ms/step
1/1 [=====] - 0s 133ms/step
1/1 [=====] - 0s 134ms/step
1/1 [=====] - 0s 134ms/step
1/1 [=====] - 0s 135ms/step
1/1 [=====] - 0s 133ms/step
1/1 [=====] - 0s 143ms/step

```

```

1/1 [=====] - 0s 155ms/step
1/1 [=====] - 0s 140ms/step
1/1 [=====] - 0s 128ms/step
1/1 [=====] - 0s 128ms/step
1/1 [=====] - 0s 132ms/step
1/1 [=====] - 0s 144ms/step
1/1 [=====] - 0s 140ms/step
1/1 [=====] - 0s 131ms/step
1/1 [=====] - 0s 130ms/step
1/1 [=====] - 0s 130ms/step
1/1 [=====] - 0s 136ms/step
1/1 [=====] - 0s 127ms/step

```

```
[75]: features = model.predict(image).reshape(2048, )
      features_list.append(features)
```

```
[76]: features_df = pd.DataFrame(features_list)
```

```
[77]: features_df["labels"] = data["labels"]
```

```
[78]: x = features_df.drop(['labels'], axis = 1)
      y = features_df.loc[:, "labels"].values
```

```
[79]: x
```

```
[79]:
```

	0	1	2	3	4	5	6	\
0	3.888965	0.003357	0.000000	0.080566	0.002496	0.003217	0.178870	
1	6.105568	0.000000	0.008155	0.063793	0.005606	0.047762	0.079004	
2	4.191666	0.000000	0.000000	0.073727	0.007410	0.040819	0.126976	
3	4.969837	0.190751	0.000000	0.060903	0.012210	0.034155	0.038725	
4	7.409759	0.000000	0.000000	0.223266	0.000000	0.013826	0.060581	
...	...	...	...	...	...	...	...	
10656	4.650623	0.000000	0.000000	0.003782	0.000000	0.000000	0.079633	
10657	5.075249	0.000000	0.000000	0.073403	0.000000	0.050068	0.109488	
10658	5.491919	0.000000	0.000000	0.062658	0.087209	0.227873	0.310514	
10659	4.076108	0.017551	0.000000	0.118688	0.000000	0.000000	0.148364	
10660	8.373823	0.000000	0.038963	0.043950	0.000000	0.000000	0.222565	
...	...	...	...	...	...	...	...	
	7	8	9	...	2038	2039	2040	\
0	0.054217	0.015593	0.044378	...	1.161147	0.046141	0.090094	
1	0.000000	0.002628	0.012909	...	0.067440	0.000000	0.000000	
2	0.026118	0.000000	0.000000	...	0.646453	0.000000	0.218996	
3	0.035803	0.183820	0.000000	...	0.090973	0.018727	0.066534	
4	0.000000	0.044893	0.025707	...	0.179018	0.005104	0.000000	
...	...	...	...	...	...	...	...	
10656	0.035964	0.052119	0.277229	...	1.089402	0.014226	0.000000	
10657	0.000000	0.127485	0.027882	...	0.697676	0.000000	0.084635	

10658	0.037401	0.000000	0.043395	...	0.154101	0.003430	0.000000
10659	0.000000	0.000000	0.000000	...	0.407699	0.026058	0.270066
10660	0.000000	0.000000	0.057938	...	0.648262	0.004149	0.000000

	2041	2042	2043	2044	2045	2046	2047
0	0.396494	0.005847	0.000000	0.605872	0.689851	0.192494	1.315786
1	0.291204	0.000000	0.000000	0.194592	0.897486	0.181279	0.888027
2	0.082670	0.000000	0.000000	0.033167	0.255844	0.011669	1.841871
3	0.355468	0.086318	0.053395	0.031325	0.539635	0.013693	1.445034
4	0.408164	0.000000	0.000000	0.100929	0.117949	0.972514	0.484715
...	...	...	...	...	...	...	...
10656	1.019124	0.584857	0.000000	0.643598	0.648637	0.098623	0.295180
10657	0.057695	0.176103	0.000000	0.666509	0.455349	1.335689	0.615851
10658	0.094444	0.133244	0.000000	0.083235	0.834814	0.120172	2.136771
10659	0.208164	0.037543	0.000000	0.560540	0.330985	0.199827	1.003238
10660	0.978938	0.364705	0.000000	0.245678	0.166369	0.471063	0.754563

[10661 rows x 2048 columns]

```
[80]: print(f"Number of features before feature selection: {x.shape[1]}")
```

Number of features before feature selection: 2048

```
[81]: y
```

```
[81]: array([1, 1, 1, ..., 0, 0, 0])
```

# Data Scaling

```
[82]: from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
scaler.fit(x)
x_ = scaler.transform(x)
```

```
[83]: x_ = pd.DataFrame(x_)
```

# Feature Selection Methods

## ANOVA

```
[84]: from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif

def anova_fs():

    selector = SelectKBest(f_classif, k=500) # k is number of features
    selector.fit(x_, y)
```

```

cols = selector.get_support(indices=True)
anova_x = x_[cols]
return anova_x

```

## Recursive Feature Elimination (RFE)

```

[85]: from sklearn.feature_selection import RFE
      from sklearn.ensemble import RandomForestClassifier

      def RFE_fs():
          rfe_selector = RFE(estimator=RandomForestClassifier())
          rfe_selector.fit(x_, y)

          rfe_support = rfe_selector.get_support()
          rfe_feature = x_.loc[:,rfe_support].columns.tolist()

          rfe_x = x_[rfe_feature]
          return rfe_x

```

## Random Forest

```

[86]: from sklearn.feature_selection import SelectFromModel
      from sklearn.ensemble import RandomForestClassifier

      def rf_fs():
          embedded_rf_selector =
          ↪SelectFromModel(RandomForestClassifier(n_estimators=200, random_state=5),
          ↪threshold='1.25*median')
          embedded_rf_selector.fit(x, y)

          embedded_rf_support = embedded_rf_selector.get_support()
          embedded_rf_feature = x.loc[:,embedded_rf_support].columns.tolist()

          rf_x = x[embedded_rf_feature]
          return rf_x

```

```

[87]: fs_x = rf_fs()

```

```

[88]: print(f"Number of features after feature selection: {fs_x.shape[1]}")

```

Number of features after feature selection: 590

# Train Test Split

```

[89]: from sklearn.model_selection import train_test_split
      x_train, x_test, y_train, y_test = train_test_split(fs_x, y, test_size = 0.2,
      ↪random_state = 42)

```

# Classification with ML Algorithms



```
[90]: from sklearn.model_selection import cross_val_score, cross_val_predict
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import GaussianNB
from sklearn.neighbors import KNeighborsClassifier
from sklearn import svm
from sklearn.metrics import confusion_matrix
from sklearn.metrics import f1_score, precision_score, recall_score, accuracy_score
from sklearn.model_selection import GridSearchCV
```

## kNN

```
[91]: neig = np.arange(1, 25)
train_accuracy = []
test_accuracy = []

for i, k in enumerate(neig):

    knn = KNeighborsClassifier(n_neighbors=k)
    knn.fit(x_train, y_train)
    prediction_ = knn.predict(x_test)
    train_accuracy.append(knn.score(x_train, y_train))
    test_accuracy.append(knn.score(x_test, y_test))

print("Best accuracy is {} with K = {}".format(np.
    ↪max(test_accuracy), 1+test_accuracy.index(np.max(test_accuracy))))
```

Best accuracy is 0.8326300984528833 with K = 15

```
[92]: knn = KNeighborsClassifier(n_neighbors=17)
knn.fit(x_train, y_train)
predicted = knn.predict(x_test)
score = knn.score(x_test, y_test)
knn_score_ = np.mean(score)

print('Accuracy : %.3f' % (knn_score_))
```

Accuracy : 0.832

```
[93]: p=precision_score(y_test, predicted)
print('Precision : %.3f' % (p))

r=recall_score(y_test, predicted)
print('Recall : %.3f' % (r))

f1=f1_score(y_test, predicted)
print('F1-score: %.3f' % (f1))

f1_w=f1_score(y_test, predicted, average='weighted')
```

```
print('Weighted f1-score: %.3f' % (f1_w))
```


Precision : 0.838

Recall : 0.935

F1-score: 0.884

Weighted f1-score: 0.825

```
[94]: cf_matrix = confusion_matrix(y_test, predicted)
sns.heatmap(cf_matrix, cmap="PuBu", annot=True, fmt='.0f')
plt.show()
```



## SVM

```
[95]: param_grid_svm = {'C': [0.1, 1, 10, 100, 1000],
                        'gamma': [1, 0.1, 0.01, 0.001, 0.0001],
                        'kernel': ['rbf', 'poly']}

SVM_grid = GridSearchCV(svm.SVC(), param_grid_svm, cv=5)
SVM_grid.fit(x_train, y_train)
```

```
[95]: GridSearchCV(cv=5, estimator=SVC(),
                  param_grid={'C': [0.1, 1, 10, 100, 1000],
```

```
'gamma': [1, 0.1, 0.01, 0.001, 0.0001],  
'kernel': ['rbf', 'poly']})
```

```
[96]: print(SVM_grid.best_params_)  
  
print(SVM_grid.best_estimator_)
```

```
{'C': 10, 'gamma': 0.01, 'kernel': 'rbf'}  
SVC(C=10, gamma=0.01)
```


```
[97]: svm_clf = svm.SVC(C=100, gamma=0.01, kernel='rbf')  
svm_clf.fit(x_train, y_train)  
predicted = svm_clf.predict(x_test)  
score = svm_clf.score(x_test, y_test)  
svm_score_ = np.mean(score)  
  
print('Accuracy : %.3f' % (svm_score_))
```

```
Accuracy : 0.899
```

```
[98]: p=precision_score(y_test, predicted)  
print('precision : %.3f' % (p))  
  
r=recall_score(y_test, predicted)  
print('recall : %.3f' % (r))  
  
f1=f1_score(y_test, predicted)  
print('f1-score: %.3f' % (f1))  
  
f1_w=f1_score(y_test, predicted, average='weighted')  
print('weighted f1-score: %.3f' % (f1_w))
```

```
precision : 0.901  
recall : 0.957  
f1-score: 0.928  
weighted f1-score: 0.897
```

```
[99]: cf_matrix = confusion_matrix(y_test, predicted)  
sns.heatmap(cf_matrix, cmap="PuBu", annot=True, fmt='.0f')  
plt.show()
```



## Random Forest

```
[100]: param_grid_rf = {
        'n_estimators': [200, 500],
        'max_depth' : [4,5,6,7,8]}

RF_grid = GridSearchCV(estimator=RandomForestClassifier(),
    param_grid=param_grid_rf, cv= 5)
RF_grid.fit(x_train, y_train)
```

```
[100]: GridSearchCV(cv=5, estimator=RandomForestClassifier(),
    param_grid={'max_depth': [4, 5, 6, 7, 8],
    'n_estimators': [200, 500]})
```

```
[101]: print(RF_grid.best_params_)
```

```
{'max_depth': 8, 'n_estimators': 500}
```

```
[102]: r_forest = RandomForestClassifier(500,max_depth=8, random_state=5)
r_forest.fit(x_train,y_train)
predicted = r_forest.predict(x_test)
score = r_forest.score(x_test, y_test)
```

```
rf_score_ = np.mean(score)

print('Accuracy : %.3f' % (rf_score_))
```

Accuracy : 0.826

```
[103]: p=precision_score(y_test, predicted)
print('precision : %.3f' % (p))


r=recall_score(y_test, predicted)
print('recall : %.3f' % (r))

f1=f1_score(y_test, predicted)
print('f1-score: %.3f' % (f1))

f1_w=f1_score(y_test, predicted, average='weighted')
print('weighted f1-score: %.3f' % (f1_w))
```

precision : 0.821  
recall : 0.952  
f1-score: 0.882  
weighted f1-score: 0.815

```
[104]: cf_matrix = confusion_matrix(y_test, predicted)
sns.heatmap(cf_matrix, cmap="PuBu", annot=True, fmt='.0f')
plt.show()
```



## Naive Bayes

```
[105]: nb_model = GaussianNB()
nb_model.fit(x_train,y_train)
predicted = nb_model.predict(x_test)
score = nb_model.score(x_test, y_test)
nb_score_ = np.mean(score)

print('Accuracy : %.3f' % (nb_score_))
```

Accuracy : 0.765

```
[106]: p=precision_score(y_test, predicted)
print('precision : %.3f' % (p))

r=recall_score(y_test, predicted)
print('recall : %.3f' % (r))

f1=f1_score(y_test, predicted)
print('f1-score: %.3f' % (f1))

f1_w=f1_score(y_test, predicted, average='weighted')
```

```
print('weighted f1-score: %.3f' % (f1_w))
```

precision : 0.849

recall : 0.797

f1-score: 0.822

weighted f1-score: 0.768

```
[107]: cf_matrix = confusion_matrix(y_test, predicted)
sns.heatmap(cf_matrix, cmap="PuBu", annot=True, fmt='.0f')
plt.show()
```

